MARKETING STRATEGY BRIEF · PART A + B

BUILDING THE IMMUNEFI MARKETING ENGINE

PAGE 1

Strategy and Priorities

SITUATION

Most startups spend years trying to earn what Immunefi already has: genuine authority in Web3 security, a community of tens of thousands of security researchers, $180B+ in assets under protection, and two products with real traction. The gap isn't credibility. It's that none of it is wired to a marketing engine yet, because there hasn't been one. That's the gap I'm here to close: building the system that finally converts a brand this strong into revenue, not starting from scratch.

The first 90 days are organized in sequence: each priority creates the conditions for the next.

PRIORITY 1: LEARN THE BUYER (WEEKS 1–4)

Before writing a single line of positioning, I need to understand why protocols actually buy. That means joining live customer conversations, interviewing 10–15 current customers across both products, and reviewing the deals we lost. Three data points matter most: what triggered the purchase, what almost killed it, and the exact language buyers use to describe the problem. The output is a validated ICP and buyer journey for each product, grounded in what customers said rather than internal assumptions.

Every piece of content, every sales enablement asset, and every community activation in this plan is only as good as the ICP it's built on. Messaging without this research reflects our assumptions about buyer language, not the buyer's own words. Errors here compound through every downstream initiative.

PRIORITY 2: ACTIVATE THE RESEARCHER COMMUNITY (WEEKS 3-8)

Trust compounds slowly; a researcher community that's dormant for another month is atrophying, not waiting. The community programs will also surface the content and signal that make Priority 3's infrastructure worth building — researchers see what Immunefi's buyers care about before anyone else does. The full activation plan is on Page 3.

PRIORITY 3: BUILD THE ENGINE (WEEKS 5-12)

With buyer truth in hand and community programs running, I build the operating layer: a messaging architecture per product, a content engine powered by the AI workflow in Part B, sales enablement for PR Reviews, and the reporting that ties marketing to pipeline. Building infrastructure before buyer insights and community programs are in place risks optimizing a system around assumptions that have not yet been validated.

Go-to-Market: Two Products, Two Motions

Security Subscriptions and PR Reviews look like one product line, but they require two different sales motions on a shared foundation.

SECURITY SUBSCRIPTIONS: A TRUST-AND-TIMING MOTION

The typical buyer, whether a founder, CTO, or security lead, already knows they need continuous coverage. The purchase fires on a trigger: a funding round, a TVL jump, a near-miss, or a competitor getting exploited. Immunefi's track record does a lot of the selling on its own: $110M+ paid out, including the $10M Wormhole bounty. My two jobs here: be impossible to miss at the trigger moment, and make the buying decision easy once they're looking.

The three tactics I would activate first:

- **Exploit post-mortems published within 48 hours** of any major Web3 incident. Not to exploit a competitor's bad day, but to be the authoritative voice exactly when buyer attention is highest. This is also the primary job for the AI pipeline in Part B.
- **Owning the search terms that fire at trigger moments:** "smart contract bug bounty platform," "managed triage service crypto," "web3 vulnerability disclosure." These are low-volume, high-intent queries. The buyer is already in-market.
- **Case studies built around payout transparency and time-to-triage benchmarks** address the two questions that kill deals late: does this actually work, and how fast? Managed Triage gets its own positioning. The MTS buyer is typically an engineering or security lead at a protocol receiving a high volume of

low-quality reports without the internal capacity to triage them accurately or quickly. The message is operational: Immunefi handles triage for roughly a third of the cost of building that function in-house, and the protocol's team acts only on what matters.

PR REVIEWS: A WORKFLOW-ADOPTION MOTION

PR Reviews embeds elite security researchers directly into a protocol's GitHub pull requests, returning findings within 24 to 72 hours. That changes who I'm selling to: the engineering lead, not the security lead. The objection shifts too, from "do we need this?" to "will this slow my devs down?" The answer is built into the product: it works inside the pull request workflow engineers already use, with no new tooling and no process change.

The three tactics I would activate first:

- **Drive qualified engineering leads to the install page.** The GitHub app install is step one of Immunefi's existing onboarding flow — the conversion mechanic is already there. Marketing's job is getting the right people to it. That means targeted placements and organic content aimed at engineering leads at protocols with active development velocity, not security teams. The message is operational: catch vulnerabilities before they reach production, without slowing the sprint.
- **Content that speaks to engineers, not security teams.** The "ship faster" and "find vulnerabilities early" angles are developer motivations. A series of short-form posts and newsletter placements showing real finding types — the class of bug PR Reviews catches that a standard code review misses — gives engineering leads proof before they commit to anything. The line I would test first: PR Reviews are the unit tests of security; audits are the final exam. It positions the product alongside the development workflow rather than as a security overhead.
- **Researcher credibility as social proof.** The researchers doing these reviews are the same people who have found and reported over 1,000 mainnet bugs and helped protect $180B in assets. That is a materially different trust signal than a generic security vendor. Short profiles, notable finds, and researcher perspectives on common PR vulnerabilities make that expertise visible to an audience that respects it.

Activating the Researcher Community

The researcher community is Immunefi's single highest-leverage marketing asset. The infrastructure is in place: All Stars, the Leaderboard, the Whitehat Hall of Fame, and Top 10 Bugs. With no marketing function in place, nobody has built the layer that turns it into demand. Four programs do that.

- **Researcher spotlight series.** Each month I identify two to three notable findings from the leaderboard or via the All Stars channel, reach out directly to the researcher with a specific ask and format (a 600 to 800 word co-bylined post, written in their voice and edited by me), and draft the piece collaboratively in Notion before the researcher reviews and approves it for publication across both their channel and Immunefi's. The outreach is personal and specific, not templated. Target cadence: first outreach in week two, first spotlight in draft by week four, two published pieces per month by day 60.

- **All Stars advocacy.** In week one I identify the five to eight All Stars who are already active on X with an audience in the protocol-security overlap. The ask: early access to new features, a co-marketing slot for their next significant find, and a simple content kit: templates they can adapt, not copy. The guiding principle is that all content uses their words, not ours. Scripted advocacy signals inauthenticity to a technically sophisticated audience and undermines the credibility this program is designed to build.
- **Rapid-response post-mortems.** Web3 exploits occur at near-weekly frequency, and Immunefi's researchers are the most qualified people alive to explain them. Researcher-reviewed analysis within 48 hours makes Immunefi the default source exactly when buyers are paying closest attention.
- **A feedback loop back into GTM.** Researchers see protocol security posture before anyone else. A monthly listening program, including AMAs, surveys, and a top-researcher council, feeds straight into product and positioning. I ran a similar program at Agoric; the listening sessions produced hundreds of insights that directly shaped GTM. Across all four programs, researchers must function as genuine partners. The value of this asset depends on that relationship remaining authentic; treating it as a distribution channel rather than a collaboration degrades it.

Token Sequencing, Dependencies, Milestones & Incident Protocol

WHERE THE TOKEN FITS

Token marketing is not a top-three priority in the first 90 days. The token launched in January 2026 and is still working through post-launch sell pressure. Pushing hard on it now, before the fundamentals are visible, spends the one thing everything else runs on: platform credibility. Tokens derive credibility from the platform beneath them; the reverse is not true.

The first 90 days are foundation only: read whether researchers are holding, selling, or ignoring the token; turn the tokenomics docs into one clear narrative with versions for researchers, protocols, and the wider ecosystem; and set a standing rule that token comms ride product milestones, never price.

That sets up a real activation push next quarter, anchored to a genuine utility launch. The clearest confirmed example is the Patrons program: community members bond $IMU behind specific researchers or protocols, share in the rewards from successful bug reports, and directly tie token utility to security outcomes. When that activation is live and producing measurable results, the narrative becomes concrete: $IMU is the mechanism through which the community backs Web3 security, not a speculative asset. That's a story about product, not price, and it survives a down market.

DEPENDENCIES

FROM PRODUCT

60–90 day roadmap visibility
Without this, I cannot sequence content around launches or set realistic Day-60 activation targets.

Named engineering contact for PR Reviews funnel
The GitHub app install is the most important marketing moment in the PR Reviews motion. Without a named contact who owns that funnel, I cannot instrument, test, or iterate it.

Product usage data: activation events, time-to-first-review, triage throughput
These are the leading indicators that tell me whether the product is landing before pipeline data is available. They are the only honest unit of account in the first 60 days.

FROM SALES

Standing access to live calls and recordings in the first 30 days
This is the primary input for the ICP and buyer journey. Everything downstream depends on it.

Shared qualified-lead definition agreed by Day 30
Without this, marketing-sourced pipeline is a number that neither function trusts.

Loss-reason data including deals that never reached a call
Churned prospects and uncontacted leads are the fastest diagnostic for the conversion blockers that no amount of top-of-funnel spend can fix.

FROM LEADERSHIP

Weekly 30-minute CEO sync for first month
I need unfiltered access to company context that isn't in any brief.

Pre-agreed crisis-communication protocol with named decision-makers
See incident protocol below. Decision authority needs to be established before it's needed.

Explicit sign-off on the Day-90 success definition
If leadership's definition is revenue and mine is leading indicators, that's a conversation I want in week one, not week thirteen.

MILESTONES

DAY 30

What I will know

Top three objections in buyer language per product. Which product has the shorter sales cycle and why. Whether the two products share a buyer. The gap between how Immunefi describes the products and how buyers describe the problem.

DAY 60

Engine running

Messaging architecture shipped for both products. PR Reviews sales enablement in field use. Two researcher spotlights published. All Stars program live with 10+ participants. Two rapid-response post-mortems shipped. Token narrative in circulation.

DAY 90

Measurable movement

Marketing-sourced and influenced pipeline reported weekly. Researcher-driven content outperforming brand-channel content on engagement. 12-month plan and hiring case presented to leadership.

No revenue commitment is attached to Day 90. With enterprise-length sales cycles, leading indicators are the honest unit of account.

A PROTOCOL ON IMMUNEFI GETS PUBLICLY EXPLOITED: MY FIRST 24 HOURS

My governing principle: get internal alignment before making an external statement, but don't let silence run past the point where it becomes its own statement.

STEP 1

> Convene a response group — CEO, security lead, client owner — in a single channel. Establish the fact base: scope and status of the exploited vulnerability, prior report history, attack status, and Safe Harbor applicability. Freeze all scheduled marketing.

STEP 2

> Draft a holding statement: awareness, contact with the affected team, update commitment. No speculation on cause, attribution, or amounts. Coordinate with the affected protocol before any substantive statement is made; the incident belongs to the protocol, and Immunefi communicates in support. Brief internal teams with an approved Q&A.

STEP 3

> Publish the holding statement if public discussion already references Immunefi; otherwise hold it against a pre-agreed trigger. Monitor researcher channels and crypto X continuously. Issue factual corrections where needed and avoid narrative engagement.

STEP 4

> Initiate the post-incident plan: a researcher-reviewed technical post-mortem, contingent on settled facts and protocol consent. Maintain a decision log for retrospective review and protocol revision.

THE HARDER SCENARIO

If the exploited vulnerability was previously reported through the Immunefi platform but not acted on by the protocol, the response changes materially. The response group expands immediately to include legal. The holding statement cannot be issued without legal review. This is not because Immunefi is at fault (responsibility for remediation sits with the protocol), but because any statement that could be read as minimizing the platform's role in disclosure must be reviewed before publication. The post-mortem is delayed until facts are fully settled and legal has cleared publication. Sustained credibility requires communicating clearly even when Immunefi's own processes intersect with the incident. That standard must hold regardless of whether Immunefi's role was procedurally correct.

Context and Reasoning

WHY THIS ORDER, NOT A FASTER ONE

The temptation as a founding hire is to pursue all initiatives simultaneously. The discipline is knowing which ones to defer. Messaging without buyer research is guessing. Community activation without messaging wastes the community's trust. Infrastructure without either just automates the mistakes. The cost is real: visible campaigns start in month two, not week two.

WHAT I AM ASSUMING, AND HOW I WOULD CHECK

01

Subscriptions demand is trigger-driven. I would check whether closed-won deals actually cluster after funding rounds and category exploits in the CRM.

02

PR Reviews friction sits with the engineering buyer. I would check the call recordings and trial funnel drop-off.

03

Researchers will co-create in public. I would check by asking ten of them in week one.

If any breaks, the plan bends: merge the messaging if the two products share one buyer; switch PR Reviews to a champion-enablement motion if friction is elsewhere; lead with infrastructure instead of community if researchers show advocacy fatigue.

WHY COMMUNITY-LED BEATS PAID-LED HERE

Web3 security buyers are technical and skeptical of advertising; trust travels through practitioners, not impressions. Immunefi already owns the biggest pool of credible practitioners in the category. Paid can amplify an engine that is already working; it cannot substitute for one. Spending on ads now would buy low-converting traffic while the asset that actually compounds sits idle.

WHAT THE EXISTING PROGRAMS TELL ME

All Stars, the Leaderboard, the Hall of Fame, and Top 10 Bugs: this is well-built community infrastructure. It rewards and ranks researchers. What's missing is the conversion layer I'd build on top: researcher stories packaged for buyers, community moments tied to product launches, advocacy with actual distribution goals.

WHAT WOULD MAKE ME CHANGE THE PLAN

BREAKING POINTS

- If buyer research shows the two products share one buyer, the dual-motion structure collapses into one.
- If researchers show advocacy fatigue, community activation moves behind infrastructure.
- If leadership's definition of Day-90 success is revenue rather than leading indicators, that is a conversation I want to have in week one, not week thirteen.

PART B

AI Workflow Design

The Security Signal-to-Content Pipeline

AI WORKFLOW

The Security Signal-to-Content Pipeline

PROBLEM

The highest-return content opportunity I see is speed and authority on security events, specifically exploits, disclosures, major payouts, and ecosystem incidents, which map directly to the Subscriptions purchase trigger in Part A. The constraint is capacity: continuously monitoring the Web3 security landscape, producing technically credible analysis, and publishing inside the 24 to 48 hour relevance window while also running the other programs in this brief is not feasible for a single marketer.

This workflow compresses a function that would otherwise require two to three staff into a review-and-publish operation for one. A similar system deployed at Agoric roughly tripled content output and cut production time by approximately 70%.

TWO APPROACHES, SAME OUTPUT

I'm building this in two forms depending on available budget and how much setup time makes sense. Both use AI at every substantive step and produce the same publish-ready output. The difference is whether the system runs automatically or whether I trigger it each morning. The human gate at the end is identical in both: no content publishes without my review, and technical claims always require researcher sign-off.

Automated pipeline ~$60-100/mo · 24/7 continuous · requires technical setup

DETECT
Continuous monitoring
Watches security feeds, researcher activity, and Immunefi platform events around the clock.
AI AUTOMATED

TRIAGE
Relevance scoring
Each signal is scored and ranked. Only high-relevance events generate a brief for review.
AI AUTOMATED

DRAFT
Content generation
Produces a full package — analysis, social thread, newsletter block — against pre-built templates.
AI AUTOMATED

REVIEW
Human gate
You + named researcher verify. Technical claims need researcher sign-off before publish.
HUMAN + AI

PUBLISH
You approve
Notion queue → blog post, X thread, newsletter block. Each approved individually.
HUMAN

Stages 4 and 5 are identical in both approaches. The human gate never leaves.

AI-assisted, human-triggered $20-40/mo flat · daily check or on alert · zero setup

DETECT

Morning scan

A pre-built prompt searches live sources and returns a ranked list of relevant security events.

HUMAN + AI

TRIAGE

You decide

Review the ranked list. Select which events cross the threshold. Takes about 3 minutes.

HUMAN

DRAFT

AI drafting

Drop the approved brief into a pre-configured AI workspace. Full package ready in 5–10 min.

HUMAN + AI

REVIEW

Human gate

You + named researcher verify. Technical claims need researcher sign-off before publish.

HUMAN + AI

PUBLISH

You approve

Notion queue → blog post, X thread, newsletter block. Each approved individually.

HUMAN

	AUTOMATED	AI-ASSISTED
Monthly cost	~$60–100 (API + hosting)	$20–40 flat
Trigger	24/7, no action needed	Daily prompt or on alert
Event → draft ready	~30–60 min	~30–60 min
Setup required	Days (technical configuration)	Minutes (zero configuration)
Output quality	High	High
Best suited for	3+ publishable events/week	First 90 days and beyond

WHAT TRIGGERS IT AND WHAT COMES OUT

A security event anywhere in Web3 — an exploit, a major payout, a significant disclosure — starts the clock. In the automated version, the system detects it immediately and begins working without any action from me. In the AI-assisted version, I run a morning scan that surfaces the most relevant events, ranked and sourced, and I select which ones to act on. Either way, the next step is the same: the system drafts a publish-ready content package against pre-built templates that encode Immunefi's voice, format, and editorial guardrails. No speculation on active incidents, no naming parties before confirmation, no unverified figures.

What arrives in the review queue, typically within 30 to 60 minutes of the event, is a complete package: a written analysis, a social thread, and a newsletter block, with sources attached. Any event involving an Immunefi client is automatically flagged and routed to the incident protocol instead. Nothing publishes without my review, and anything involving technical security claims requires sign-off from a named researcher before it goes out.

HOW I WOULD KNOW IT IS WORKING: FOUR METRICS

01 · Speed

Time from event to publish

Tracked from the date of the triggering event to the date of publication.

02 · Compression

Marketer hours per published piece

I log time spent in review and editing, not in drafting. The gap between the two is the system's contribution.

03 · Quality

Drafts approved with minor edits

A rising approval rate means the templates are calibrating well. A declining rate means the templates need recalibration, not more volume.

04 · Business impact

Event-driven vs. evergreen traffic

If rapid-response content does not outperform evergreen by Day 60, I reconsider what events are worth covering, not how fast to cover them.

WHY THIS ONE, FIRST

Plenty of marketing work can be automated. This is the piece worth building before anything else, regardless of which version you start with. Both approaches make Immunefi the fastest credible voice on the events that already move its buyers. The AI-assisted version is the right starting point for the first 90 days: zero infrastructure cost, zero setup time, same output quality. The automated version earns its complexity only when event volume exceeds what a morning scan can manage, typically three or more publishable events per week. Either way, it frees me up for the work no model can do: the customer interviews, the researcher relationships, and the judgment calls that require a human with context in the room.

One tension worth naming: this workflow serves the Subscriptions trigger, not PR Reviews, which the assignment identifies as a significant and growing priority. That is a deliberate choice. Exploit post-mortems build category authority that benefits both products, and PR Reviews is built for self-serve acquisition through the GitHub install. The content pipeline creates the awareness that warms both. A dedicated PR Reviews content program is the logical next build once messaging architecture from Priority 1 is in place.